

# Wicked Fast PaaS

## Performance Tuning of OpenShift v3 and Docker

# Environment Setup

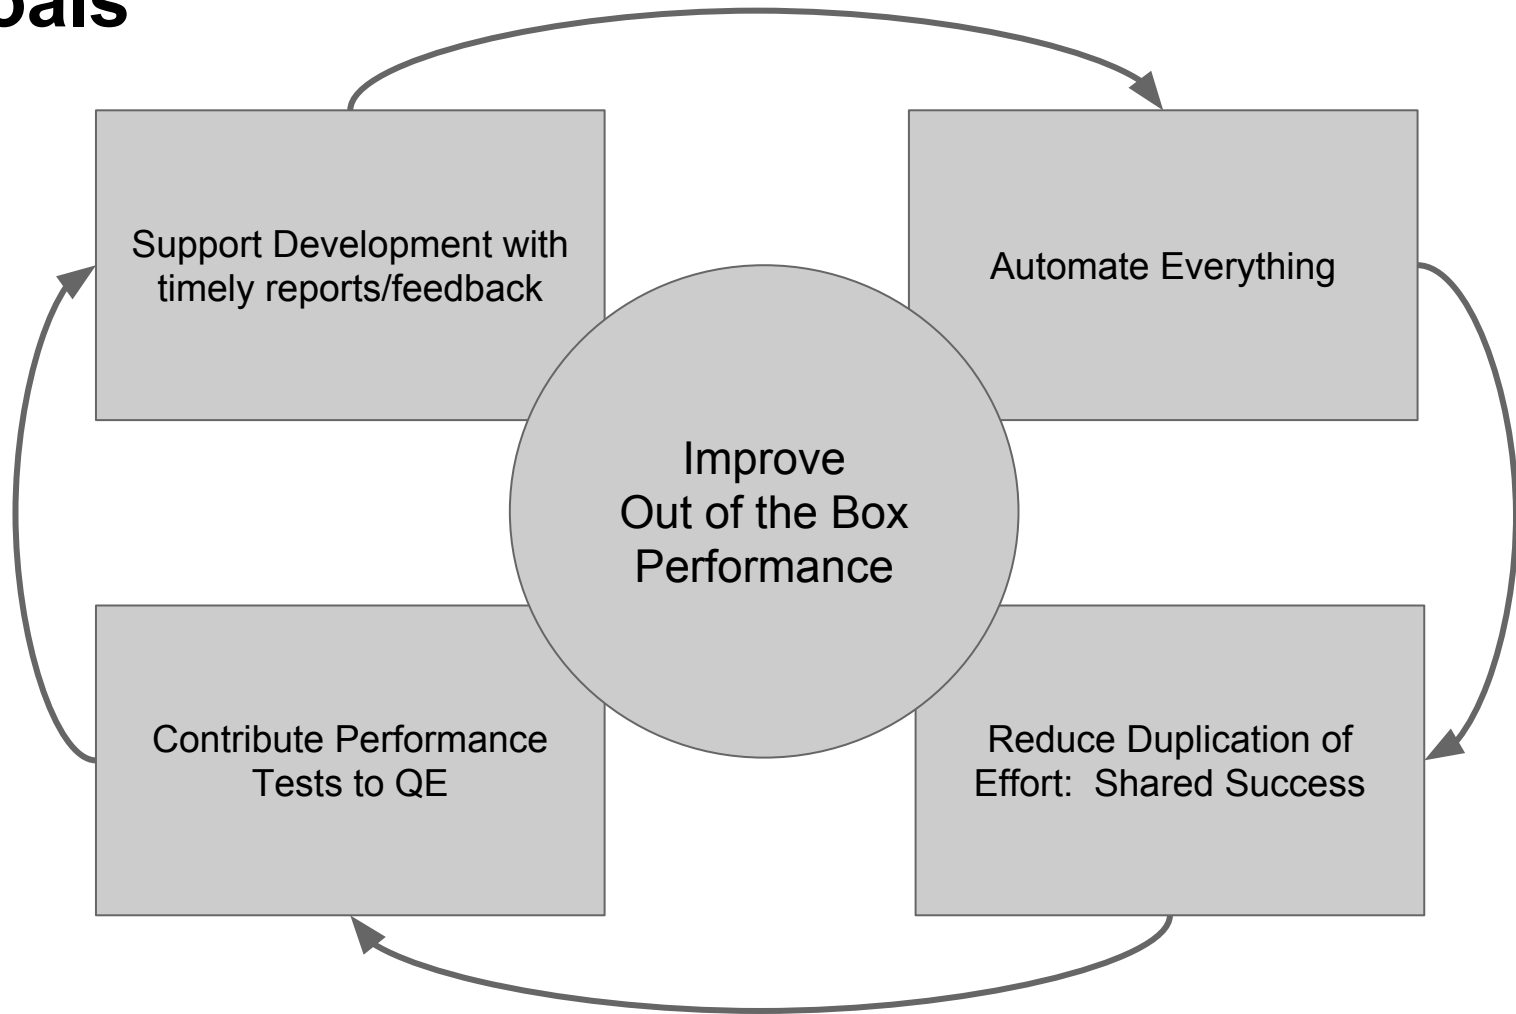
<https://github.com/jeremyeder/openshift-performance>

- Download OVA file (or copy from USB disk)
- Install VirtualBox and kernel-devel RPM that matches your running kernel
- `systemctl restart system-modules-load` or Reboot to load kernel modules
- Start VirtualBox
- Go to File -> Import Appliance -> Select the OVA file
- Click the checkbox to reset the MAC address
- Click Import and Start the VM
- Username: devconf2016 Password: devconf2016
- User devconf2016 has sudo access

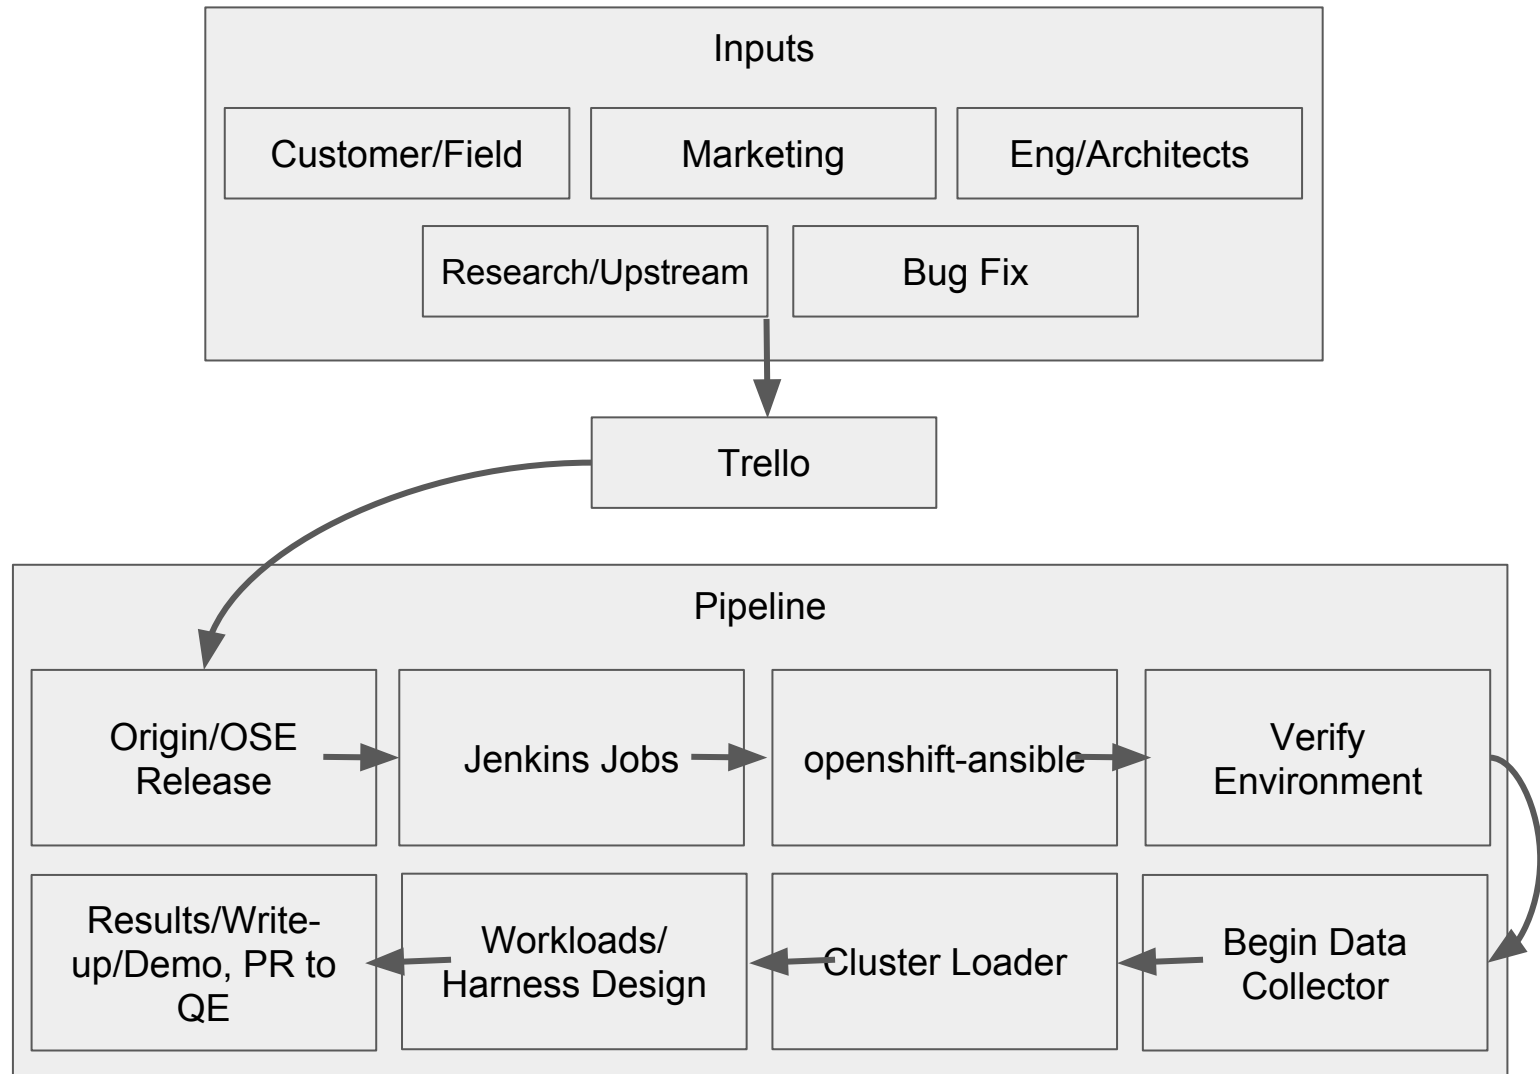
# Agenda

- Approach to Performance Analysis
- Latest Features
- Infrastructure Optimization
  - Compute, Network, Storage
- Tuning Docker and OpenShift
- Scaling OpenShift
- Architecture Overview

# Goals



# Workflow



# Tuning/Scaling fundamentals

**don't change  
just for containers**

# First: Tuning the Installer :-)

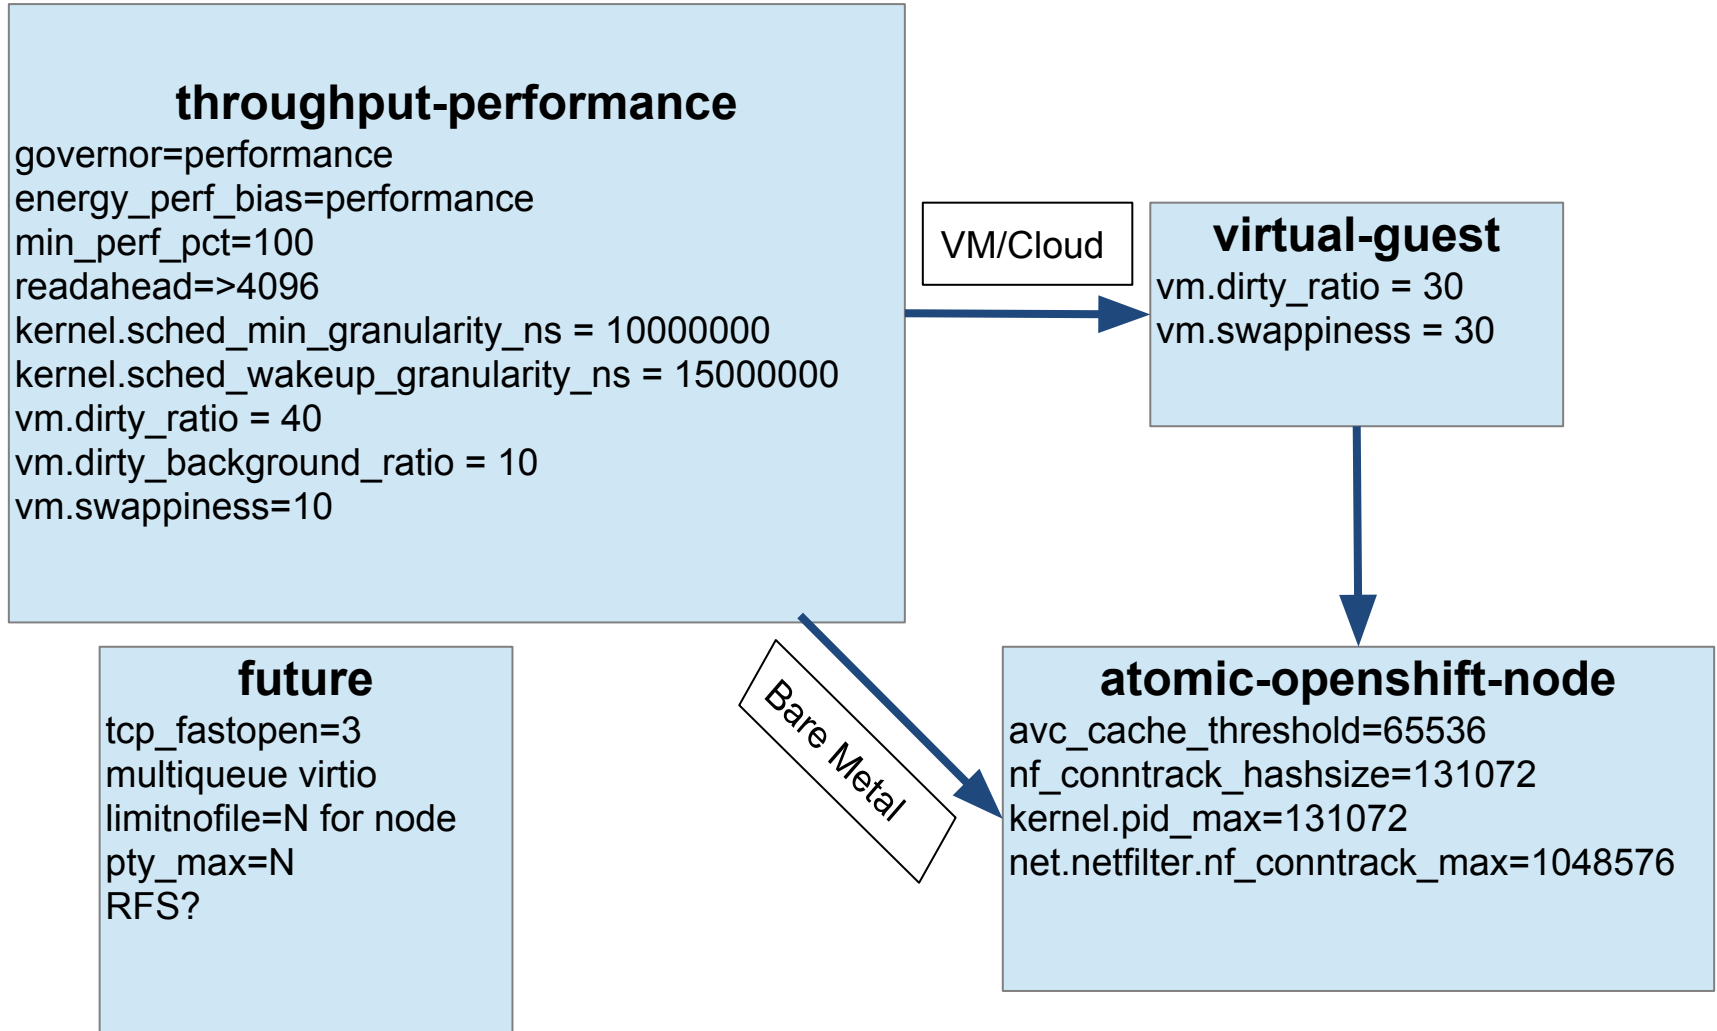
- Parallelizing is good and bad
  - Creates high load on content source.
- Installer node should be RHEL6.6 or later (ControlPersist)
- Pre-seed everything possible into your “gold image”
  - OS Updates, docker, docker-storage-setup, docker images, pre-register with Satellite/Content Source
- Ensure fast-access to content: Red Hat CDN/Satellite
- Ansible
  - Set forks  $\geq$  nodes
  - Installer should be run on same LAN as cluster.
  - Increase ControlPersist to maintain persistent SSH connection

# Ansible Config for Large Clusters

```
[defaults]
forks = 1000
host_key_checking = False
remote_user = root
roles_path = roles/
gathering = smart
fact_caching = jsonfile
fact_caching_connection = /tmp/$USER_ansible/facts
fact_caching_timeout = 600
log_path = /tmp/$USER_ansible.log
[privilege_escalation]
become = False
[ssh_connection]
ssh_args = -o ControlMaster=auto -o ControlPersist=600s
control_path = %(directory)s/%%h-%%r
pipelining = True
```



# Tuned Profiles for OpenShift



# Pbench

A Framework for Benchmarking and  
Performance Analysis

<https://github.com/distributed-system-analysis/pbench>

# What is Pbench?

- pbench (perf bench) aims to:
  - provide easy access to benchmarking & performance tools on Linux systems
  - standardize the collection of telemetry and configuration information
  - automate benchmark execution
  - output effective visualization for analysis
  - allow for ingestion into elastic search

# rhel-tools container

rhel-tools (and fedora-tools and centos-tools) are purpose-built analysis and debugging “super privileged containers”.

- strace, tcpdump, sysstat, sosreport, git

[Overview](#) and [Official Documentation](#)

tl;dr

```
# docker pull centos/tools
```

```
# atomic run centos/tools
```

# Resource Management

```
$ cat openshift-performance/svt/content/quota-default.json
  "memory": "1Gi", # every pod can use 1Gi of memory
  "cpu": "20", # "milli-cores"
  "pods": "10", # max pods
  "services": "5", # max services
  "replicationcontrollers":"5", # max rc's
  "resourcequotas":"1" # max quota objects
```

<https://github.com/kubernetes/kubernetes/blob/master/docs/design/resources.md>

# CPU/Memory Optimization

- RHEL7 task scheduler adds automatic numa\_balancing
- Pod commands can use numactl
  - docker has support for --cpuset-cpus and --cpuset-mems
  - Not in Kube yet
  - Pod manifests can use nodeSelector w/node labels to land on fast gear

# Storage Optimization

- *Ensure you are using thinLVM (not loopLVM)*
- Persistent data gets stored in “Persistent Volumes”
  - Ceph/Gluster/NFS/iSCSI/Fiber
  - Bind-mounted into container at startup
- Container storage I/O plays by the same rules as always
- I/O scheduler and others (vm.dirty etc) are system-wide
- If a container has a very particular tuning need, consider dedicated resources (HostPath pass-through)

# Docker Graph Driver

- Pluggable image/container storage backend
- Device Mapper
  - Use [docker-storage-setup](#), which will setup “thinLVM”
  - Supported in RHEL7.0+, SELinux and POSIX compliant
- Overlay FS
  - Supported (with important caveats) as of RHEL7.2
  - Increased density, faster container start/stop (page cache sharing)
  - Non-POSIX compliant, no SELinux support
- Comparison <https://developerblog.redhat.com/2014/09/30/overview-storage-scalability-docker/>



# Network Optimization


- OpenShift and Atomic Enterprise “just need connectivity”
- We use OpenvSwitch w/VXLAN tunnels
- VXLAN handles 1G pipes w/o issue
  - 10G+ needs tuning: VXLAN-offload, faster CPUs, jumbo frames
- Container network I/O plays by the same rules as always
  - NIC-level tuning such as jumbo frames/offloads are interface-wide
  - Kernel sysctl tunings are often per-container (somaxconn), sometimes not (tcp\_mem)
  - You can share host network stack or use kernel-bypass into a container (Solarflare OpenOnload/Intel DPDK)
  - <http://developerblog.redhat.com/2015/04/09/accelerating-rhel7-linux-containers-solarflare-openonload/>
  - <http://developerblog.redhat.com/2015/06/02/can-you-run-intels-data-center-network-stack-in-a-container/>

# VXLAN-offload

- Certain NICs have VXLAN-offload capabilities in hardware
  - High-end models from Intel, Mellanox, Emulex, and RHEL7.1+
- VXLAN-offload handles packet checksums on the NIC rather than the CPU
- Another in a long line of hardware-assist (MMX, SSE, AVX, GRO, TSO...)
- Public clouds not offering this yet

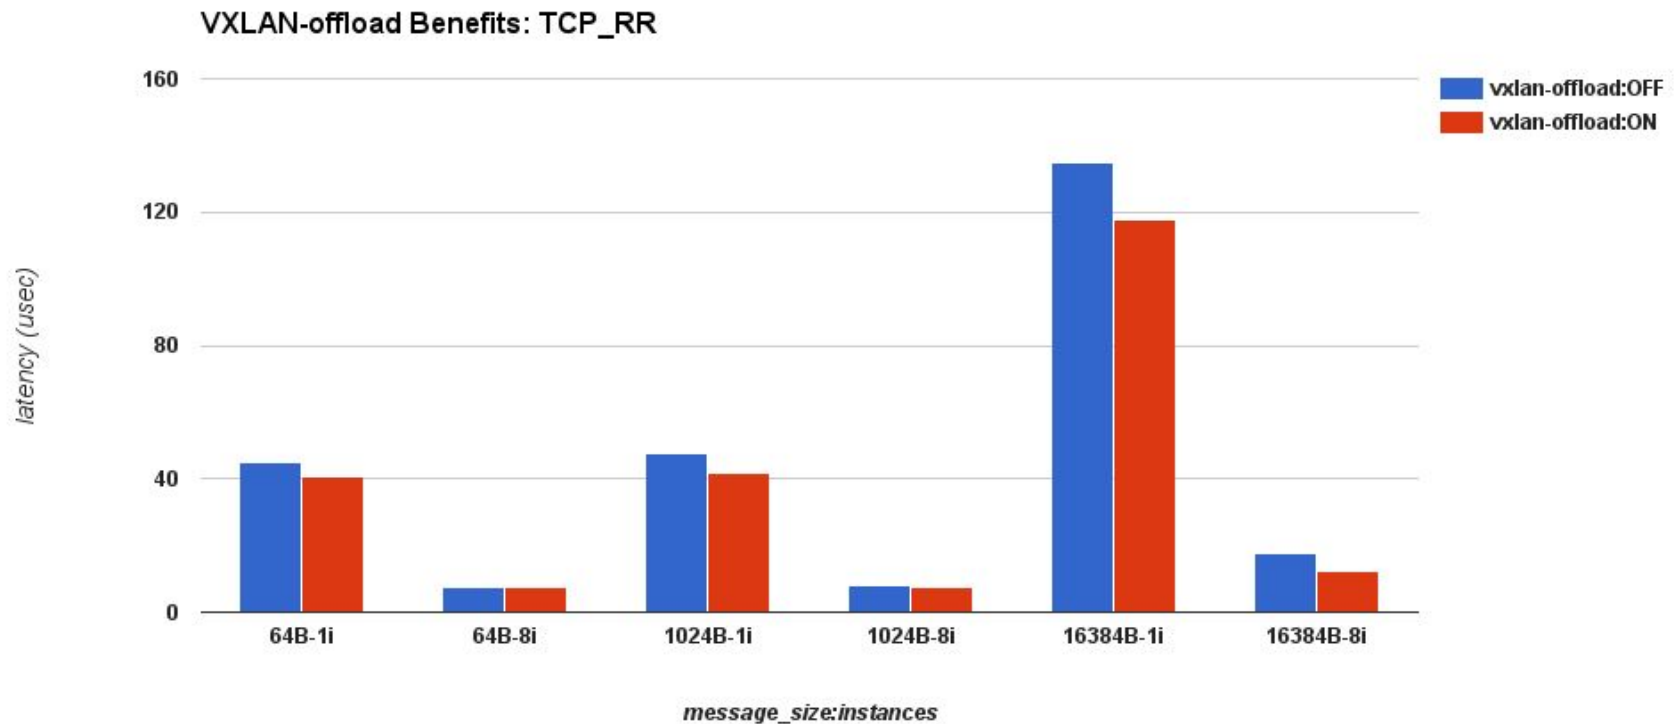
Bare Metal

# Benefits of VXLAN-offload (throughput)

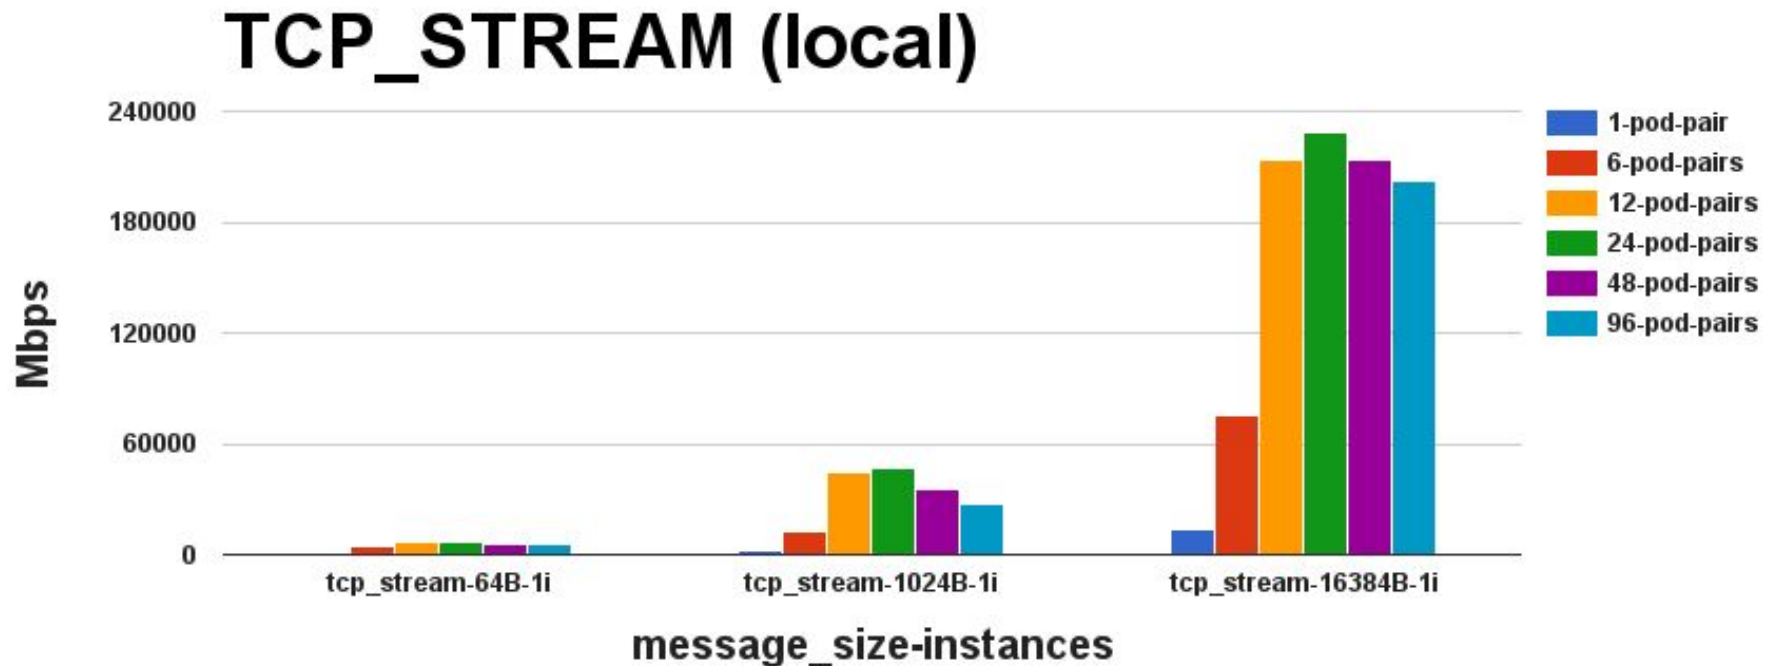


Bare Metal

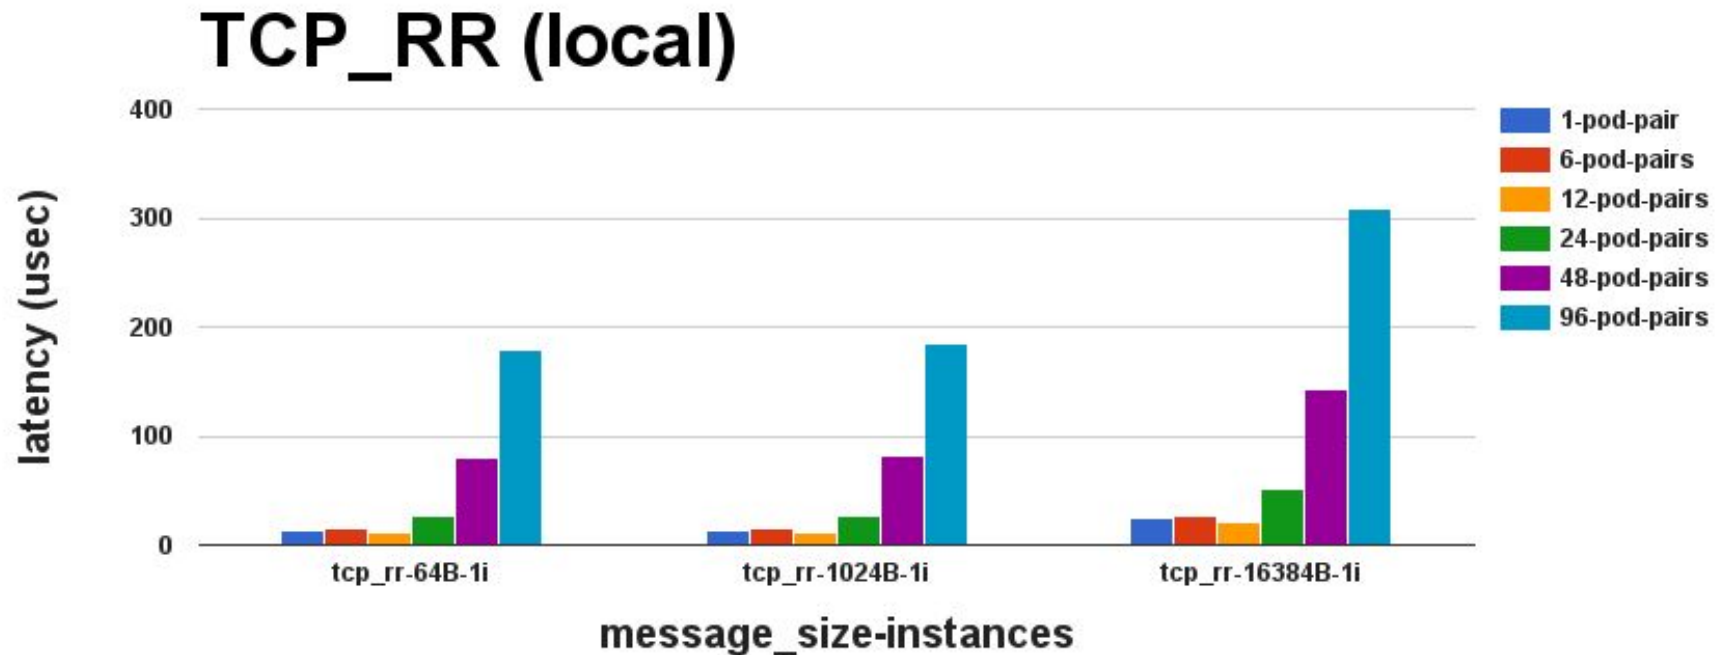
# Benefits of VXLAN-offload (latency)



# Network Performance (on-box: many pods)



# Network Performance (on-box: many pods)



# Node Heartbeat Optimization

- Cluster network communication shares media with Pod traffic
  - Extreme network load can block cluster heartbeats and lead to node eviction
  - Increase `--node-monitor-grace-period` in `/etc/origin/node/node-config.yml`


```
apiServerArguments: null
controllerArguments:
  node-monitor-grace-period:
    - "120s"
```

# Cloud Gotchas...

- Variable performance
- Pay-for-performance
- Reset your expectations




# Gotcha: EBS I/O Credits/Bursting



*Depending ~~on~~ I/O rate,  
this may never recover*

# Gotcha: EBS I/O Credits/Bursting

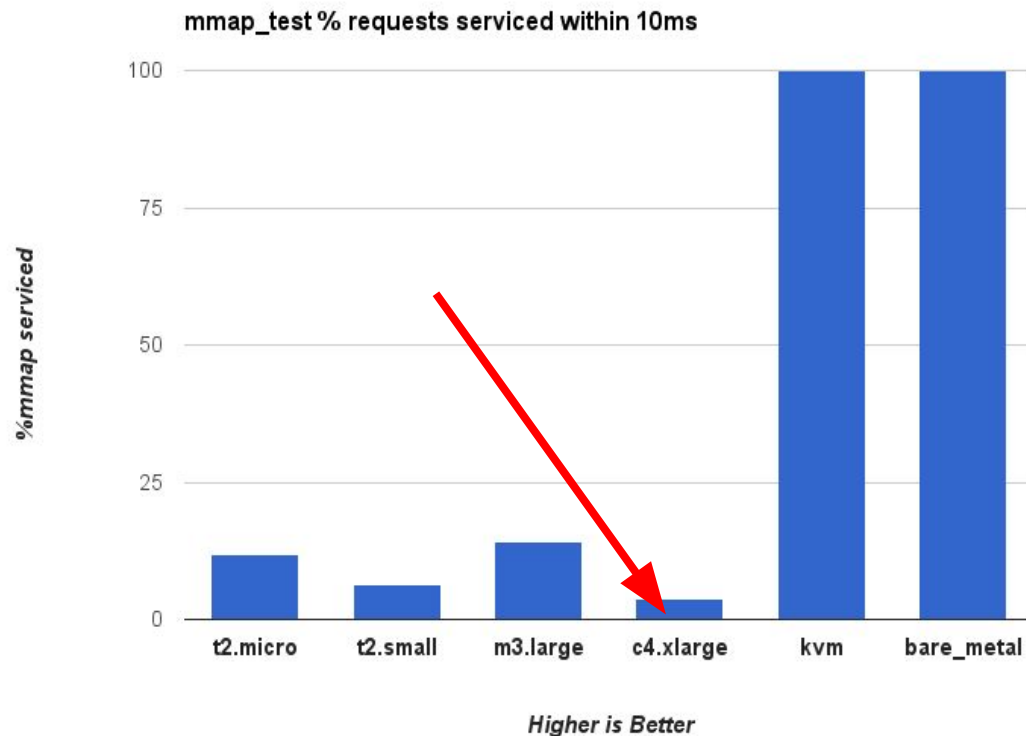


*Depending on I/O rate,  
this may never recover*


# Gotcha: CPU Credit System

*Pay for  
performance?*

*No...pay for  
determinism and  
stability*



# Storage Performance



# Out-of-the-Box Limits

- Kubelet limits to 40 pods per node by default
- OpenShift SDN: 255 nodes, 255 IPs per node
  - Tunable
- Device Mapper backend does not permit page-cache sharing
- Master node does not host pods by default
- Individual “projects” can have Resource Quotas

```
kubeletArguments:
```

```
  max-pods:
```

```
    - "NN"
```

# etcd Tuning

- Needs fast disk (SSD preferred)
- Uses RAM for snapshots...efficiency and performance improvements coming
  - snapshot efficiencies
  - reduction in garbage collection pauses
- [https://github.com/coreos/etcd/tree/master/Documentation/benchmark\\_s](https://github.com/coreos/etcd/tree/master/Documentation/benchmark_s)
- Avoid swap
- Optimize connection between etcd and master
  - Or co-locate them on the same machine

# Kubernetes Pod Manifests

- Recipe for how to deploy an application
- Some tuning options are exposed through Kube
- Encapsulate tuning inside script.sh (numactl)

pseudo-code manifest:

- image: my-app:1.0
- securityContext:
  - privileged: false
  - capabilities:
    - add:
      - CAP\_SYS\_ADMIN
- command:
  - /your/script.sh
- volumeMounts:
  - mountPath: "/perf1"

# Scheduler Options

/etc/origin/master/scheduler.json

**MatchNodeSelector** # land a pod on certain nodes

**PodFitsResources** # ensure sufficient resources to run a pod

**PodFitsPorts** # ensure ports are available

**NoDiskConflict** # enforce single writer

**Region serviceAffinity labels=region** # keep services in same region

**LeastRequestedPriority weight: 1** # favor less-committed nodes

**ServiceSpreadingPriority weight: 1** # spread pods between nodes

**Zone", "weight" : 2, serviceAntiAffinity label=zone** # keep service on different zones



# Profiling OpenShift (golang)

- Append `OPENSIFT_PROFILE={cpu,mem,web}` to `/etc/sysconfig/openshift-master`

```
# systemctl restart openshift-master
```

```
# systemctl stop openshift-master
```

```
# go tool pprof /bin/openshift /var/lib/openshift/cpu.pprof
```

```
# top10 -cum
```

- Example: <https://github.com/openshift/origin/issues/5106>

<https://github.com/openshift/origin/blob/master/HACKING.md#performance-debugging>

# **@ Scale Deployment matters**


# Scaling Up $O(10^2+)$

HA is important

- master-api & master-controller becomes a single point of failure and has caps to prevent system overload.
  - Load balancing (master-api)
  - active-passive on master-controller
  - in-flight-request limit: 400

# Kubernetes Resync Interval

- Identified and fixed a periodic load spike that limited scale
- [Issue #5106](#)



*Spike reduced in both frequency and magnitude*

# OpenShift v3 HA, Scale-out Architecture on EC2

